

# Numeric Prediction, Logistic Regression, and Feature Selection

Blake Pappas

2023-12-17

## Part I. Numeric Prediction in R

```
library(caret)
library(rpart)
library(rpart.plot)
library(glmnet)
library(dplyr)
```

Import “purchase.csv” into R:

```
purchase = read.csv("purchase.csv")
```

1. Use 5-fold cross-validation to evaluate the performance of the regression tree model on this dataset. Report the average (1) MAE; (2) MAPE; (3) RMSE.

Then, build a single regression tree model on the entire dataset, plot the tree, and answer the following two questions.

- (a) How many decision nodes are there in the tree?
- (b) Pick one decision rule from the tree and interpret it.

```
# Cross-Validation
cv = createFolds(y = purchase$Spending, k = 5)

tree_mae_cv = c()

tree_mape_cv = c()
```

```

tree_rmse_cv = c()

for (test_row in cv) {

  purchase_train = purchase[-test_row, ]
  purchase_test = purchase[test_row, ]

  tree = rpart(Spending ~ ., data = purchase_train)

  pred_tree = predict(tree, purchase_test)

  # MAE
  tree_mae = mean(abs(pred_tree - purchase_test[, 23]))
  mae_cv = c(tree_mae_cv, tree_mae)

  # MAPE
  tree_mape = mean(abs((pred_tree - purchase_test[, 23]) / purchase_test[, 23]))
  mape_cv = c(tree_mape_cv, tree_mape)

  # RMSE
  tree_rmse = sqrt(mean((pred_tree - purchase_test[, 23])^2))

  rmse_cv = c(tree_rmse_cv, tree_rmse)
}

# Average MAE
print(mean(mae_cv))

## [1] 101.1279

#Average MAPE
print(mean(mape_cv))

## [1] 1.404391

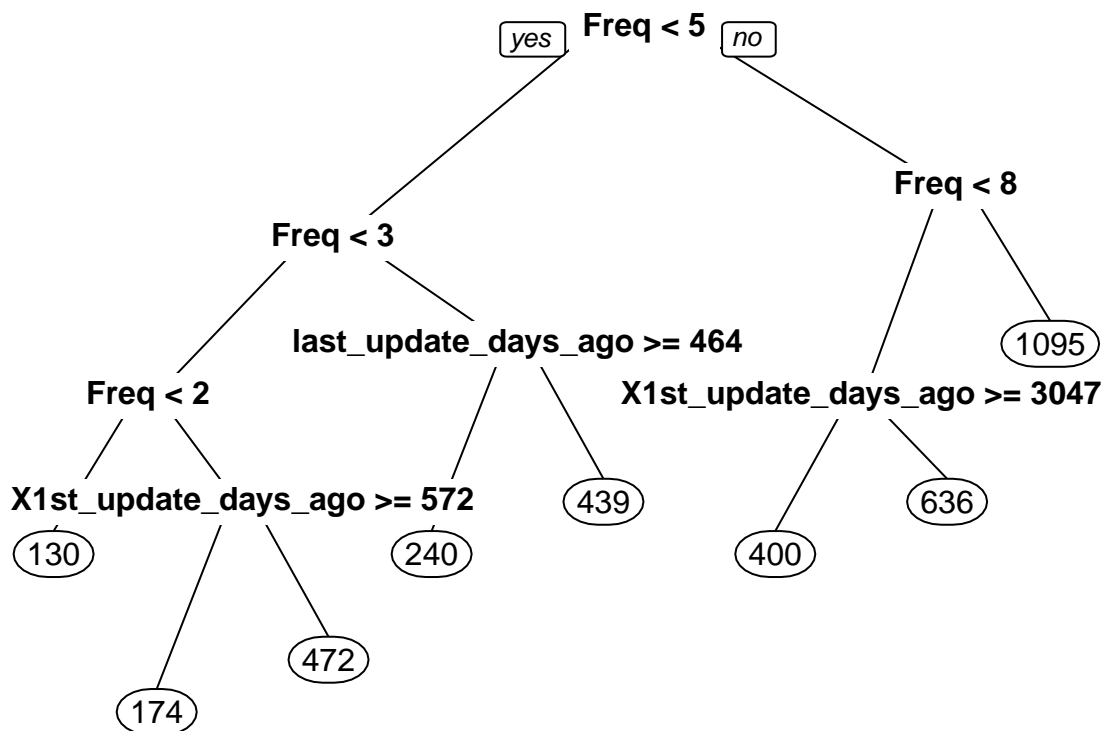
# Average RMSE
print(mean(rmse_cv))

## [1] 173.8571

# Build a Single Regression Tree Model
tree = rpart(Spending ~ ., data = purchase_train)

# Plot the Tree
prp(tree, varlen = 0)

```



Answer: There are \_\_\_\_ decision nodes in the tree.

Answer:

2. Use 5-fold cross-validation to evaluate the performance of the linear regression model on this dataset. Report the average (1) MAE; (2) MAPE; (3) RMSE.

Then, build a single linear regression model on the entire dataset and examine the model.

Pick a coefficient and interpret it.

```

# Cross-Validation
cv = createFolds(y = purchase$Spending, k = 5)

lm_mae_cv = c()

```

```

lm_mape_cv = c()

lm_rmse_cv = c()

for (test_row in cv) {

  purchase_train = purchase[-test_row, ]
  purchase_test = purchase[test_row, ]

  lm_model = lm(Spending ~ ., data = purchase_train)

  pred_lm = predict(lm_model, purchase_test)

  # MAE
  lm_mae = mean(abs(pred_lm - purchase_test[, 23]))
  mae_cv = c(lm_mae_cv, lm_mae)

  # MAPE
  lm_mape = mean(abs((pred_lm - purchase_test[, 23]) / purchase_test[, 23]))
  mape_cv = c(lm_mape_cv, lm_mape)

  # RMSE
  lm_rmse = sqrt(mean((pred_lm - purchase_test[, 23])^2))

  rmse_cv = c(lm_rmse_cv, lm_rmse)
}

# Average MAE
print(mean(mae_cv))

```

```
## [1] 90.39615
```

```

#Average MAPE
print(mean(mape_cv))

```

```
## [1] 1.106005
```

```

# Average RMSE
print(mean(rmse_cv))

```

```
## [1] 148.9209
```

```

# Build a Single Linear Regression Model
lm_model = lm(Spending ~ ., data = purchase_train)

# Examine the Model
summary(lm_model)

```

```

##
## Call:
## lm(formula = Spending ~ ., data = purchase_train)

```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -469.43  -90.36  -16.23   50.54 1201.50
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    7.297e+01  6.059e+01   1.204   0.229
## US             -1.558e+01  1.660e+01  -0.939   0.348
## source_a        5.206e+01  5.830e+01   0.893   0.372
## source_c       -1.777e+01  6.346e+01  -0.280   0.780
## source_b        2.156e+01  6.436e+01   0.335   0.738
## source_d       -7.009e+00  6.319e+01  -0.111   0.912
## source_e        5.760e+00  5.895e+01   0.098   0.922
## source_m      -1.995e+01  7.070e+01  -0.282   0.778
## source_o        1.407e+01  7.972e+01   0.176   0.860
## source_h      -1.361e+02  8.327e+01  -1.634   0.103
## source_r        8.840e+01  6.045e+01   1.462   0.144
## source_s        2.693e+00  6.660e+01   0.040   0.968
## source_t        8.370e-01  6.678e+01   0.013   0.990
## source_u        3.827e+01  5.835e+01   0.656   0.512
## source_p        7.945e+00  7.779e+01   0.102   0.919
## source_x      -4.509e+00  7.235e+01  -0.062   0.950
## source_w        7.802e+00  5.946e+01   0.131   0.896
## Freq           9.492e+01  5.135e+00 18.486 < 2e-16 ***
## last_update_days_ago -9.784e-03  1.025e-02  -0.955   0.340
## X1st_update_days_ago -1.684e-02  1.207e-02  -1.395   0.163
## Web.order        2.476e+00  1.201e+01   0.206   0.837
## Gender.male       6.726e+00  1.202e+01   0.559   0.576
## Address_is_res   -8.065e+01  1.470e+01  -5.487 5.54e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 167.3 on 777 degrees of freedom
## Multiple R-squared:  0.4965, Adjusted R-squared:  0.4822
## F-statistic: 34.83 on 22 and 777 DF,  p-value: < 2.2e-16
```

A one-unit increase in X is associated with b-units increase in Prediction Y.

## Part II. Logistic Regression and Feature Selection

```
library(FSelectorRcpp)
library(pROC)
```

1. Import the data. Convert the “spam” variable to a factor.

```
spambase = read.csv("spambase.csv")  
spambase$spam = factor(spambase$spam)
```

2. Use 5-fold cross-validation to evaluate the performance of the logistic regression model on this dataset. Report the average (1) accuracy; (2) precision, recall, and F-measure of class “spam”; (3) AUC of class “spam”.

```
cv = createFolds(y = spambase$spam, k = 5)  
accuracy = c()  
  
for (test_row in cv) {  
  spambase_train = spambase[-test_row, ]  
  spambase_test = spambase[test_row, ]  
  
  logit_model = glm(spam ~ ., data = spambase_train,  
                    family = binomial(link = "logit"))  
  
  # Log Odds  
  pred = predict(logit_model, spambase_test)  
  
  # Predicted Probability  
  pred_prob = predict(logit_model, spambase_test,  
                     type = "response")  
  
  # Binary Predictions  
  pred_binary = ifelse(pred_prob > 0.5, "yes", "no")  
  
  cm = confusionMatrix(factor(pred_binary), factor(ifelse(spambase_test$spam == "1", "yes", "no")),  
                      mode = "prec_recall", positive = "yes")  
  
  accuracy = c(accuracy, cm$overall[1])  
}  
  
# Average Accuracy  
print(mean(accuracy))  
  
## [1] 0.9271864  
  
# Precision  
print(cm$byClass[5])
```

```
## Precision
## 0.9088398
```

```
# Recall
print(cm$byClass[6])
```

```
## Recall
## 0.9063361
```

```
# F-Measure
print(cm$byClass[7])
```


```
## F1
## 0.9075862
```

```
# AUC
roc_logit = roc(response = spambase_test$spam,
                 predictor = pred_prob)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_logit)
```



```
auc(roc_logit)
```

```
## Area under the curve: 0.9724
```

**3. Perform feature selection using the information gain metric. Build the best logistic regression with the selected features.**

**Note:** Use a for loop to try different numbers of features. The “best” model is defined as the one with highest AUC for class “spam”.

**Report the features and AUC of the “best” model.**

```
best_auc = 0
best_model = vector()

for (i in 1:57) {
  IG = information_gain(spam ~ ., data = spambase_train)

  topK = cut_attr(IG, k = i)

  train = spambase_train %>% select(topK, spam)
  test = spambase_test %>% select(topK, spam)

  logit_model = glm(spam ~ ., data = train, family = binomial(link = "logit"))
  pred_prob = predict(logit_model, test, type = "response")
  auc = auc(ifelse(test$spam == 1, 1, 0), pred_prob)
  if (auc > best_auc){
    best_auc = auc
    best_model = topK
  }
}
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```



[illegible]

[illegible]

```
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1
```

[illegible]

```
## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1
```

```
best_model
```

```
## [1] "char_freq..3" "char_freq..4"
## [3] "capital_run_length_longest" "word_freq_remove"
## [5] "word_freq_your" "capital_run_length_average"
## [7] "word_freq_free" "word_freq_money"
## [9] "capital_run_length_total" "word_freq_000"
## [11] "word_freq_hp" "word_freq_george"
## [13] "word_freq_you" "word_freq_our"
## [15] "word_freq_hpl" "word_freq_receive"
## [17] "word_freq_all" "word_freq_business"
## [19] "word_freq_address" "word_freq_credit"
## [21] "word_freq_internet" "word_freq_mail"
## [23] "word_freq_order" "word_freq_over"
## [25] "word_freq_email" "word_freq_will"
## [27] "char_freq..5" "word_freq_addresses"
## [29] "char_freq..1" "word_freq_re"
## [31] "word_freq_1999" "word_freq_lab"
## [33] "word_freq_labs" "word_freq_make"
## [35] "word_freq_85" "word_freq_edu"
## [37] "word_freq_650" "word_freq_meeting"
## [39] "word_freq_telnet" "word_freq_people"
## [41] "word_freq_857" "word_freq_original"
## [43] "word_freq_415" "word_freq_pm"
## [45] "word_freq_report" "word_freq_data"
## [47] "word_freq_project" "word_freq_cs"
## [49] "word_freq_technology" "word_freq_conference"
## [51] "char_freq..2" "word_freq_font"
## [53] "word_freq_3d" "char_freq."
## [55] "word_freq_direct" "word_freq_table"
## [57] "word_freq_parts"
```

```
best_auc
```

```
## Area under the curve: 0.9724
```

4. Perform forward feature selection. Build the best logistic regression with the selected features.

The “best” model is defined as the one with highest AUC for class “spam”, evaluated using 5-fold cross-validation.

Report the features and AUC of the “best” model.

### Forward Feature Selection

```
best_auc = 0
selected_features = c()
while (TRUE) {
  feature_to_add = -1
  # The elements of setdiff(x, y) are those elements in x, but not in y
  for (i in setdiff(1:57, selected_features)) {
    train = spambase_train %>% select(selected_features, i, spam)
    test = spambase_test %>% select(selected_features, i, spam)
    spambase_best = spambase %>% select(selected_features, i, spam)

    logit_model = glm(spam ~ ., data = train,
                      family = binomial(link = "logit"))

    pred_prob = predict(logit_model, test,
                        type = "response")

    auc = auc(test$spam, pred_prob)

    if (auc > best_auc) {
      best_auc = auc
      feature_to_add = i
    }
  }

  if (feature_to_add != -1) {
    selected_features = c(selected_features, feature_to_add)
  }
  else break
}

print(selected_features)
```

```
## [1] 52 7 53 25 27 46 5 42 44 9 20 8 16 24 41 12 33 39 26 48 10 35 29
```

```
print(best_auc)
```

```
## Area under the curve: 0.9786
```

## Cross-Validation

```
cv = createFolds(y = spambase_best$spam, k = 5)

accuracy = c()

for (test_row in cv) {

  spambase_best_train = train
  spambase_best_test = test

  logit_model = glm(spam ~ ., data = spambase_best_train,
                    family = binomial(link = "logit"))

  # Log Odds
  pred = predict(logit_model, spambase_best_test)

  # Predicted Probability
  pred_prob = predict(logit_model, spambase_best_test,
                     type = "response")

  # Binary Predictions
  pred_binary = ifelse(pred_prob > 0.5, "yes", "no")

  cm = confusionMatrix(factor(pred_binary), factor(ifelse(spambase_best_test$spam == "1", "yes", "no")))
  mode = "prec_recall", positive = "yes")

  accuracy = c(accuracy, cm$overall[1])
}

# Confusion Matrix
print(cm)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no yes
##           no 534 39
##           yes 23 324
##
##           Accuracy : 0.9326
##           95% CI : (0.9144, 0.9479)
##           No Information Rate : 0.6054
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.8579
```


```
##
## McNemar's Test P-Value : 0.05678
##
##           Precision : 0.9337
##           Recall    : 0.8926
##           F1        : 0.9127
##           Prevalence : 0.3946
##           Detection Rate : 0.3522
##           Detection Prevalence : 0.3772
##           Balanced Accuracy : 0.9256
##
##           'Positive' Class : yes
##
```

```
# AUC
roc_logit = roc(response = spambase_best_test$spam,
                 predictor = pred_prob)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_logit)
```





```
auc(roc_logit)
```

```
## Area under the curve: 0.977
```

## Performance Evaluation

```
# Confusion Matrix  
print(cm)
```


```
## Confusion Matrix and Statistics  
##  
##           Reference  
## Prediction  no  yes  
##           no  534  39  
##           yes   23 324  
##  
##           Accuracy : 0.9326  
##           95% CI : (0.9144, 0.9479)  
##           No Information Rate : 0.6054  
##           P-Value [Acc > NIR] : < 2e-16  
##  
##           Kappa : 0.8579  
##  
##           McNemar's Test P-Value : 0.05678  
##  
##           Precision : 0.9337  
##           Recall : 0.8926  
##           F1 : 0.9127  
##           Prevalence : 0.3946  
##           Detection Rate : 0.3522  
##           Detection Prevalence : 0.3772  
##           Balanced Accuracy : 0.9256  
##  
##           'Positive' Class : yes  
##
```

```
# AUC  
roc_logit = roc(response = spambase_best_test$spam,  
                predictor = pred_prob)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_logit)
```



```
auc(roc_logit)
```

```
## Area under the curve: 0.977
```