

Completely Randomized Designs - Lab

Blake Pappas

December 17, 2023

The dataset `PlantGrowth` contains results from an experiment to compare yields (as measured by dried weight of plants) obtained under a control and two different treatment conditions.

Data Source: Dobson, A. J. (1983) *An Introduction to Statistical Modelling*. London: Chapman and Hall.

Let's load the data first:

Code:

```
data(PlantGrowth)
str(PlantGrowth)
```

```
## 'data.frame': 30 obs. of 2 variables:
## $ weight: num 4.17 5.58 5.18 6.11 4.5 4.61 5.17 4.53 5.33 5.14 ...
## $ group : Factor w/ 3 levels "ctrl","trt1",...: 1 1 1 1 1 1 1 1 1 1 ...
```

1. Compute the treatment means and standard deviations.

Code:

```
(means <- tapply(PlantGrowth$weight, PlantGrowth$group, mean))
```

```
## ctrl trt1 trt2
## 5.032 4.661 5.526
```


```
(vars <- tapply(PlantGrowth$weight, PlantGrowth$group, sd))
```

```
## ctrl trt1 trt2
## 0.5830914 0.7936757 0.4425733
```


2. Make side-by-side boxplots by treatment.

Code:

```
boxplot(weight ~ group, data = PlantGrowth, las = 1)
```



```
boxplot(weight ~ group, data = PlantGrowth, las = 1, horizontal = T)
```



3. Write down the effects model and explain each term in the model (including the model assumptions regarding the random error).

Answer: $Y_{ij} : \mu + \alpha_i + \epsilon_{ij}$, $i = 1, \dots, g$, $j = 1, \dots, n_i$, $\epsilon_{ij} \sim N(0, \sigma^2)$. Y_{ij} is the random variable that represents the response for the j th experimental unit to treatment i . μ_i is the mean response time for the i th treatment. ϵ_{ij} is the random error, which is assumed to be normal.

4. Perform an overall F-test using ANOVA. State the hypotheses, p-value, decision, and conclusion.

Code:

```
AOV <- aov(weight ~ group, data = PlantGrowth)
summary(AOV)
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## group      2  3.766   1.8832   4.846 0.0159 *
## Residuals 27 10.492   0.3886
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Answer: The null hypothesis is $H_0 : \alpha_i = 0$, the alternative hypothesis is $H_a : \alpha_i \neq 0$, the α is 0.05, the F-statistic is 4.846, the p-value is 0.0159, and the decision is to reject the null hypothesis. The data shows evidence of differences the three treatment groups.

5. Conduct pairwise comparisons using Tukey's HSD procedure.

Code:

```
# HSD
HSD <- TukeyHSD(AOV, conf.level = 0.95)
HSD$group

##           diff          lwr          upr      p adj
## trt1-ctrl -0.371 -1.0622161 0.3202161 0.39087114
## trt2-ctrl  0.494 -0.1972161 1.1852161 0.19799599
## trt2-trt1  0.865  0.1737839 1.5562161 0.01200642
```


Answer: The trt2-trt1 groups have statistically significant differences.

6. Use qqplot to examine the normality assumption on error.

Code:

```
mod1 <- lm(weight ~ group, data = PlantGrowth)

qqnorm(mod1$resid, cex = 0.8)
qqline(mod1$resid, col = "red", lwd = 1.5)
```




Answer: Based on the plot, the modeling assumption of normality appears to be reasonable. The pattern of the plot runs very close to the trend line.


7. Make a residual plot to assess equal variance assumption.

Code:

```
plot(mod1$fitted, mod1$resid, las = 1, xlab = "Fitted Value", ylab = "Residual", cex = 0.75, col = "blue",  
abline(h = 0))
```




```
plot(mod1$resid, col = "blue", las = 1, cex = 0.75, pch = 16, ylab = "Residual")  
abline(h = 0, lty = 2, lwd = 1.5)
```



```
# By Treatment
attach(PlantGrowth)
colorf <- as.factor(PlantGrowth$group)
colors <- c("green",
"orange", "blue")

plot(mod1$resid, col = colors[colorf], las = 1, cex = 0.75,
      pch = 16, ylab = "Residual")
abline(h = 0, lty = 2, lwd = 1.5)
```



Answer: Based on the residual plots, the assumption of equal variances seems to be reasonable.

8. Perform a Levene's test for equal variance.

Code:

```
# Levene's Test for Equal Variance
# install.packages("lawstat")
library("lawstat")
levene.test(PlantGrowth$weight, PlantGrowth$group, location = "mean")

##
## Classical Levene's test based on the absolute deviations from the mean
## ( none not applied because the location is not set to median )
##
## data: PlantGrowth$weight
## Test Statistic = 1.237, p-value = 0.3062

# Brown-Forsythe Test
levene.test(PlantGrowth$weight, PlantGrowth$group, location = "median")

##
## Modified robust Brown-Forsythe Levene-type test based on the absolute
## deviations from the median
##
```

```
## data:  PlantGrowth$weight  
## Test Statistic = 1.1192, p-value = 0.3412
```